

```

! pip install kaggle

! mkdir ~/.kaggle
! cp kaggle.json ~/.kaggle/
! chmod 600 ~/.kaggle/kaggle.json

! kaggle datasets download romanleo2003/labtinkoff
! unzip -O utf8 "/content/labtinkoff.zip"

import torch
import torch.nn as nn
import numpy as np
import torchvision
import torchvision.transforms as transforms
import torch.nn.functional as F
from torchvision import datasets, models
from PIL import Image
import os
import torch.utils.data as data_utils
import matplotlib.pyplot as plt

DATASET_PATH = "../content/CCPD2019-dl1"
dir = os.path.abspath(os.curdir)
data_dir=os.path.join(dir, "CCPD2019-dl1")
data_dir

{"type": "string"}

```

Создаем класс датасета как наследника torch.utils.data.Dataset

```

class GetData(torch.utils.data.Dataset):

    def __init__(self, path, transform = None, train = True):

        self.transform = transform
        if train:
            self.path = os.path.join(path, 'train')
        else:
            self.path = os.path.join(path, 'test')
        self.listdir = os.listdir(self.path)

    def __len__(self):
        return len(self.listdir)

    def __getitem__(self, x):
        if torch.is_tensor(x):
            x = x.tolist()

```

```

    img_name = os.path.join(self.path, self.listdir[x])
    image = Image.open(img_name)
    if self.transform:
        image = self.transform(image)
    label = img_name[-10:-4]
    data_item = {'image': image, 'label': label}
    return data_item

batch_size = 64
car_num_transform =
transforms.Compose([transforms.ToTensor(), transforms.Resize((85, 256))])
)
train_data = GetData(data_dir, train = True, transform =
car_num_transform)
test_data = GetData(data_dir, train = False, transform =
car_num_transform)
train_loader = torch.utils.data.DataLoader(dataset = train_data,
                                             batch_size = batch_size)
test_loader = torch.utils.data.DataLoader(dataset = test_data,
                                         batch_size = batch_size)

```

Здесь реализованы основные функции, которые нам понадобятся:
разбитие картинок и меток на символы, перевод букв в метки классов (у
нас их будет 34: 10 цифр и (26 - 2) буквы, так как тут не будет символов '0' и
'T' (я вроде узнал, что в китайских номерах их нет, что и логично: их трудно
отличить от '0' и '1' соответственно, и буквы с цифрами стоят почти в
разнобой)

```

mask = [35, 75, 115, 147, 180, 212]

def split_to_literals(x):
    literals = list(torch.tensor_split(x, mask, axis = x.dim() - 1))
    literals.pop(0)
    return literals

def split_labels(x):
    size = len(x)
    list_labels = [torch.zeros(size) for x in range(6)]
    for i in range(size):
        splitted_label = list(x[i])
        for j in range(6):
            list_labels[j][i] = class_to_number(splitted_label[j])
    return list_labels

def class_to_number(x):
    ind = ord(x)
    if(ind >= 48 and ind <= 57):
        return ind - ord('0')
    elif(ind < 73):

```

```

        return ind - ord('A') + 10
    elif(ind > 73 and ind < 79):
        return ind - ord('A') + 9
    else:
        return ind - ord('A') + 8

def number_to_class(x):
    if(x <= 9):
        return chr(ord('0') + x)
    elif(x <= 17):
        return chr(ord('A') + x - 10)
    elif(x <= 22):
        return chr(ord('A') + x - 9)
    else:
        return chr(ord('A') + x - 8)

```

Смотрим на данные

```

plt.figure(figsize=(25, 25))
print('===== TRAIN =====')
for i in range(8):
    image = train_data[i]
    ax = plt.subplot(1, 8, i + 1)
    ax.set_title(image['label'])
    ax.axis('off')
    plt.imshow(image['image'].permute(1, 2, 0))

plt.show()
plt.figure(figsize=(25, 25))
print('===== TEST =====')
for i in range(8):
    image = test_data[i]
    ax = plt.subplot(1, 8, i + 1)
    ax.set_title(image['label'])
    ax.axis('off')
    plt.imshow(image['image'].permute(1, 2, 0))
plt.show()

```

===== TRAIN =====



===== TEST =====



Смотря на данные, можно понять:

1. Число символов фиксированно
2. Номера хорошо просегментированы и одинаково спроектированы

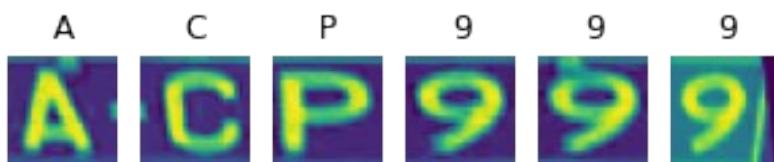
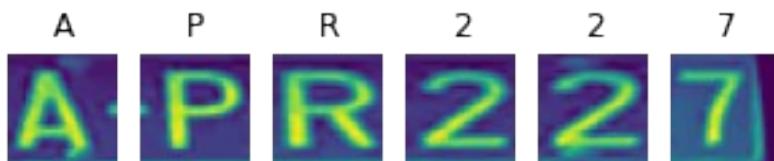
3. Нет особо никакой последовательной взаимосвязи между символами (то есть не как в российских - где-то буквы, где-то цифры (только буква вначале, но ее грех не распознать, тк она более удалена от остальных 5-ти))

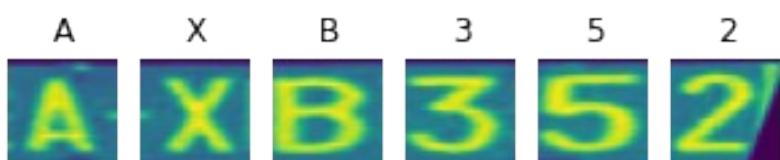
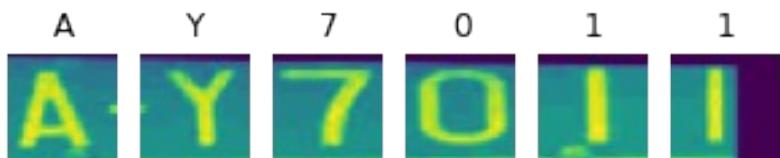
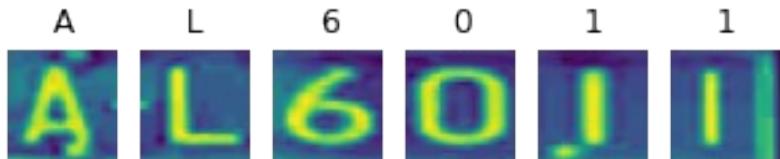
Поэтому:

1. В рекурентных слоях нет необходимости из-за отсутствия последовательной связи
2. Однаковая распределенность позволяет разбить номер на символы, регионы расположения одинаковы для всех номеров
3. Отдельные символы распознавать можно и с простой CNN + residual слой, просто нужно сделать эффективное по вычислениям разбиение на символы в процессе подгрузки данных с учетом батчей
4. Конец

```
transform_liters =
transforms.Compose([transforms.Resize((85, 75)), transforms.CenterCrop((75, 75)),
transforms.Resize((28, 28)), torchvision.transforms.Normalize((0.1307),
(0.3081)), transforms.Grayscale(num_output_channels=1)])
```

```
for i in range(5):
    K = train_data[i]
    liters = split_to_liters(K['image'])
    for i in range(len(liters)):
        ax = plt.subplot(1, len(liters) + 1, i + 1)
        ax.set_title(K['label'][i])
        ax.axis('off')
        plt.imshow(transform_liters(liters[i]).squeeze(0))
plt.show()
```





Как видно все норм разбивается

```
class Res(nn.Module):
    def __init__(self,
                 in_channels: int,
                 out_channels: int
                 ):
        super().__init__()
        self.in_channels = in_channels
        self.out_channels = out_channels
        self.ds = torch.nn.Conv2d(
            in_channels,
            out_channels,
            kernel_size=1
        )
        self.conv1 = torch.nn.Conv2d(
            in_channels,
            out_channels,
            kernel_size=3,
            padding=1
        )
        self.conv2 = torch.nn.Conv2d(
            out_channels,
            out_channels,
            kernel_size=3,
            padding=1
        )
        self.relu = torch.nn.ReLU(inplace=True)
    def forward(self, input: torch.Tensor) -> torch.Tensor:
        buff = input
```

```
x = self.conv1(input)
x = self.relu(x)
x = self.conv2(x)

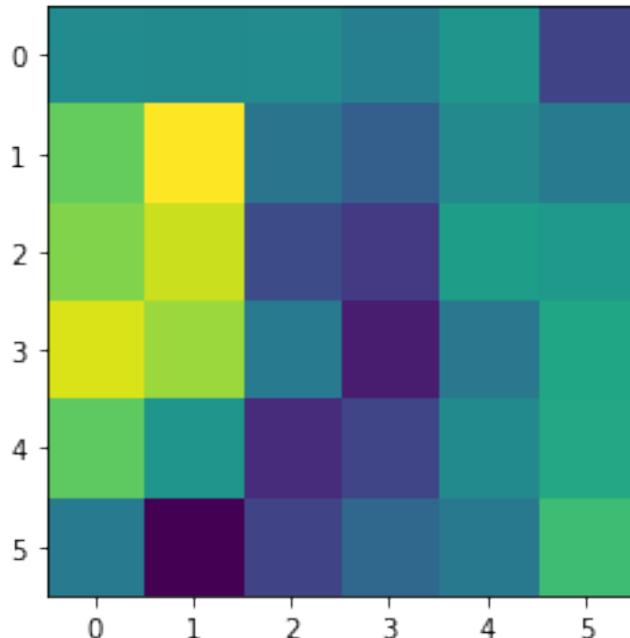
if (self.in_channels != self.out_channels):
    buff = self.ds(buff)

x += buff
x = self.relu(x)

return x

j = cnn_block = torch.nn.Sequential(
    torch.nn.Conv2d(1, 16, 7),
    torch.nn.ReLU(),
    torch.nn.Conv2d(16, 32, 5),
    torch.nn.ReLU(),
    torch.nn.Conv2d(32, 64, 5),
    torch.nn.ReLU(),
    torch.nn.Conv2d(64, 64, 3),
    torch.nn.ReLU(),
    Res(64,64),
    torch.nn.MaxPool2d(2)
)
plt.figure()
K = train_data[0]
liters = split_to_liters(K['image'])
J = j(transform_liters(liters[0]))
plt.imshow(J[7].detach())
print(J.shape)
plt.show()

torch.Size([64, 6, 6])
```



```

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.cnn_block = torch.nn.Sequential(
            torch.nn.Conv2d(1, 16, 7),
            torch.nn.ReLU(),
            torch.nn.Conv2d(16, 32, 5),
            torch.nn.ReLU(),
            Res(32,32),
            torch.nn.Conv2d(32, 64, 5),
            torch.nn.ReLU(),
            torch.nn.Conv2d(64, 64, 3),
            torch.nn.ReLU(),
            Res(64,64),
            torch.nn.MaxPool2d(2)
        )
        self.clf = torch.nn.Linear(6 * 6 * 64, 24 + 10)

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        feature_map = self.cnn_block(x)
        feature_vector = torch.flatten(feature_map, x.dim() - 3)
        return self.clf(feature_vector)

net = Net()
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
optimizer = torch.optim.SGD(net.parameters(), lr=0.05)
criterion = torch.nn.CrossEntropyLoss()
net.to(device);

```

```

from tqdm import tqdm
from typing import Callable, NamedTuple, List
from collections import namedtuple

def train_epoch(
    model: torch.nn.Module,
    optimizer: torch.optim.Optimizer,
    loader: torch.utils.data.DataLoader,
    criterion: torch.nn.modules.loss._Loss,
    device: torch.device,
    verbose: bool = False,
) -> float:
    model.train(True)
    model.to(device)
    optimizer.zero_grad()
    acc_loss = 0
    total = len(loader.dataset)
    if verbose:
        loader = tqdm(loader, desc="Training", total=len(loader),
    leave=True)
    for input_data in loader:
        liters = split_to_liters(input_data['image'])
        labels = split_labels(input_data['label'])
        for i in range(6):
            input = transform_liters(liters[i]).to(device)
            target = labels[i].type(torch.LongTensor)
            target = target.to(device)

            predicted = model(input)

            loss = criterion(predicted, target)

            # calculate gradient
            loss.backward()
            # update weights
            optimizer.step()
            # flush gradients
            optimizer.zero_grad()
            acc_loss += loss.item()

    return acc_loss / total

```

```
EvalOutput = namedtuple("EvalOutput", ["loss", "accuracy"])
```

```

def eval_epoch(
    model: torch.nn.Module,
    loader: torch.utils.data.DataLoader,
    criterion: torch.nn.modules.loss._Loss,

```

```

        device: torch.device,
        verbose: bool = False,
    ) -> EvalOutput:
        model.train(False)
        model.to(device)
        acc_loss = 0
        acc = 0
        total = len(loader.dataset)
        # no grad for context manager to accelerate evaluation
        with torch.no_grad():
            if verbose:
                loader = tqdm(loader, desc="Evaluation",
total=len(loader), leave=True)
            for input_data in loader:
                liters = split_to_liters(input_data['image'])
                labels = split_labels(input_data['label'])
                for i in range(6):
                    input = transform_liters(liters[i]).to(device)
                    target = labels[i].type(torch.LongTensor)
                    target = target.to(device)
                    predicted = model(input)

                    loss = criterion(predicted, target)
                    acc_loss += loss.item()

                    acc += torch.sum(
                        torch.argmax(predicted, 1) == target
                    ).item()

            return EvalOutput(
                loss=acc_loss / total / 6,
                accuracy=acc / total / 6
            )

    class TrainOutput(NamedTuple):
        train_loss: List[float]
        val_loss: List[float]
        val_accuracy: List[float]

    def train(
        num_epochs: int,
        model: torch.nn.Module,
        optimizer: torch.optim.Optimizer,
        train_loader: torch.utils.data.DataLoader,
        test_loader: torch.utils.data.DataLoader,
        criterion: torch.nn.modules.loss._Loss,
        device: torch.device
    ) -> TrainOutput:

```

```

train_loss = []
val_loss = []
val_acc = []
for epoch in range(num_epochs):
    loss = train_epoch(
        model, optimizer, train_loader, criterion, device,
        verbose=True
    )
    train_loss.append(loss)
    eval_out = eval_epoch(
        model, test_loader, criterion, device, verbose=True
    )
    val_loss.append(eval_out.loss)
    val_acc.append(eval_out.accuracy)

    print(f"Epoch #{epoch}:")
    print(f"Training Loss: {loss}")
    print(f"Evaluation Loss: {eval_out.loss}")
    print(f"Accuracy: {eval_out.accuracy}")

return TrainOutput(
    train_loss=train_loss,
    val_loss=val_loss,
    val_accuracy=val_acc
)

```

num_epochs = 5

```

training_results = train(
    num_epochs,
    net,
    optimizer,
    train_loader,
    test_loader,
    criterion,
    device
)

```

Training: 100%|██████████| 3125/3125 [09:43<00:00, 5.36it/s]
Evaluation: 100%|██████████| 157/157 [00:26<00:00, 6.03it/s]

Epoch #0:
Training Loss: 0.0034646165605452216
Evaluation Loss: 0.0007442258928709243
Accuracy: 0.9888322165549889

Training: 100%|██████████| 3125/3125 [09:24<00:00, 5.53it/s]
Evaluation: 100%|██████████| 157/157 [00:21<00:00, 7.31it/s]

Epoch #1:
Training Loss: 0.00010849459648283088

```

Evaluation Loss: 0.0005389873720008297
Accuracy: 0.9912324565789912

Training: 100%|██████████| 3125/3125 [09:01<00:00, 5.78it/s]
Evaluation: 100%|██████████| 157/157 [00:21<00:00, 7.28it/s]

Epoch #2:
Training Loss: 7.139462926044459e-05
Evaluation Loss: 0.0005036425654328657
Accuracy: 0.9920658732539921

Training: 100%|██████████| 3125/3125 [08:52<00:00, 5.87it/s]
Evaluation: 100%|██████████| 157/157 [00:20<00:00, 7.58it/s]

Epoch #3:
Training Loss: 5.460775743324978e-05
Evaluation Loss: 0.00058047389253935
Accuracy: 0.9912324565789912

Training: 100%|██████████| 3125/3125 [08:44<00:00, 5.96it/s]
Evaluation: 100%|██████████| 157/157 [00:20<00:00, 7.53it/s]

Epoch #4:
Training Loss: 4.320976056281075e-05
Evaluation Loss: 0.0004951160394569068
Accuracy: 0.9923158982564924

```

```

def predict(img):
    liters_list = split_to_liters(img)
    word = ''
    for i in range(6):
        pred = net(transform_liters(liters_list[i]).to(device))
        x = torch.argmax(pred)
        word = word + number_to_class(x)
    return word

with torch.no_grad():
    for i in range(100):
        K = test_data[i]
        plt.figure()
        plt.imshow(K['image'].permute(1,2,0))
        plt.title(predict(K['image']))
        plt.show()

```

AAN098



AMC799



AM136W



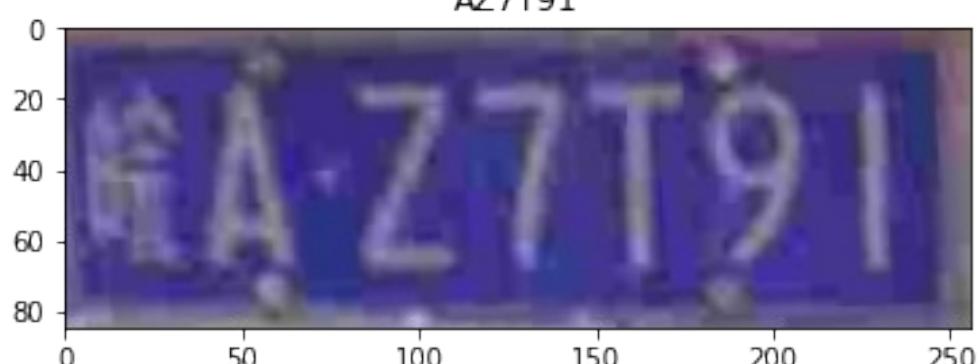
NOF338



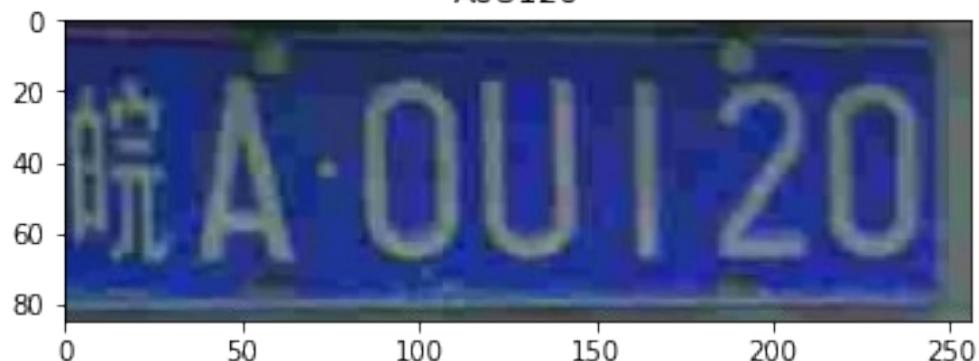
AXA253



AZ7T91



AOU120



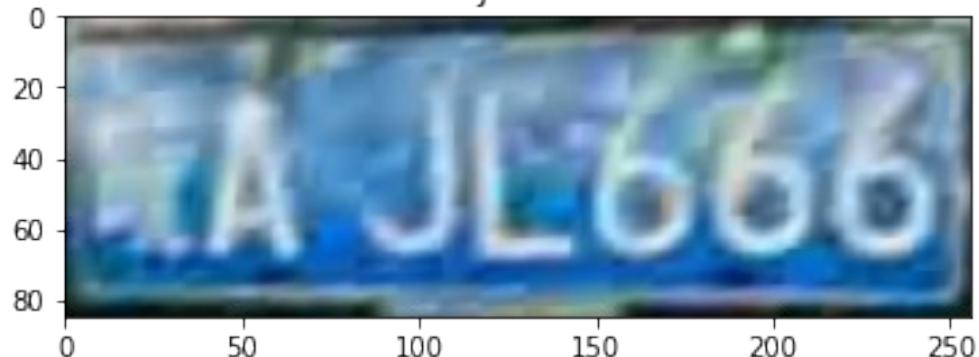
AW2S13



AZS493



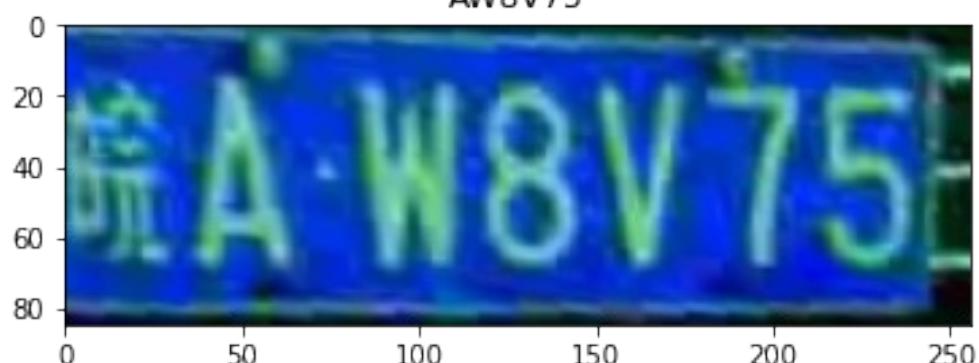
AJE666



AS4982



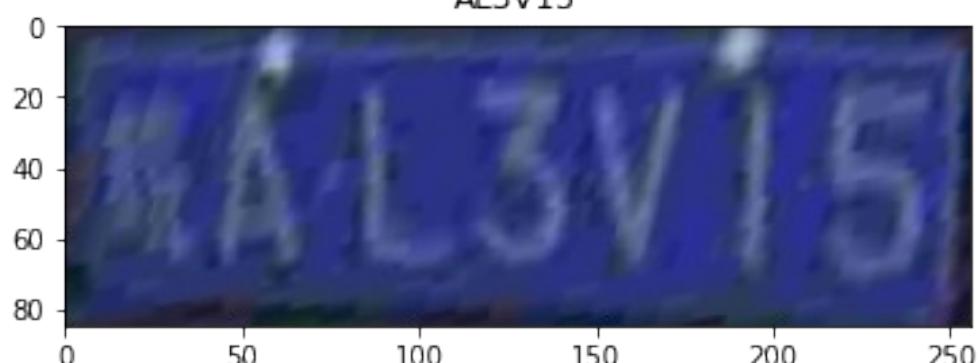
AW8V75



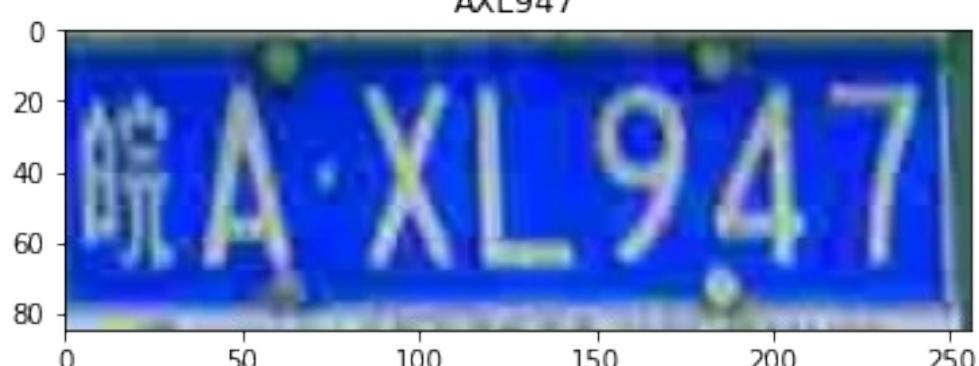
AL7900



AL3V15



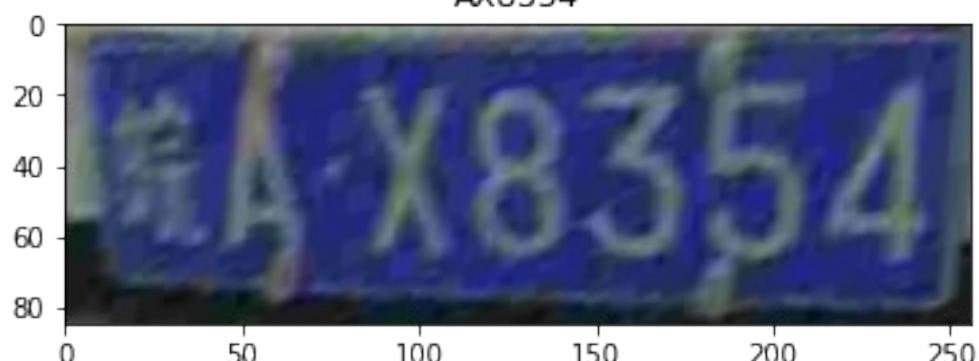
AXL947



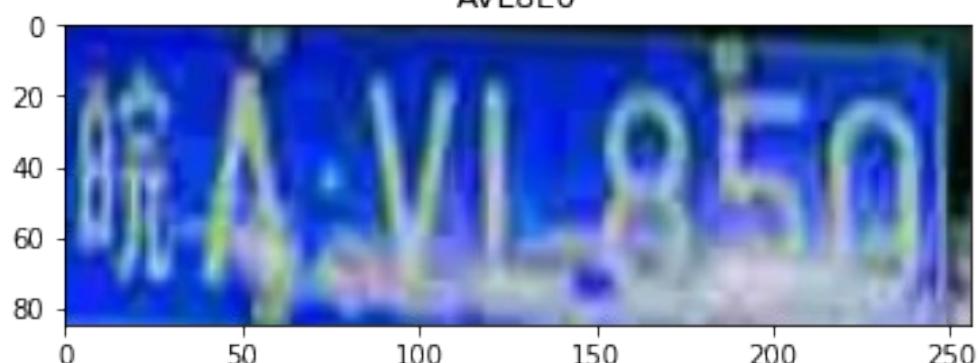
ACM3F9



AX8354



AVL8E0



A18RQ0



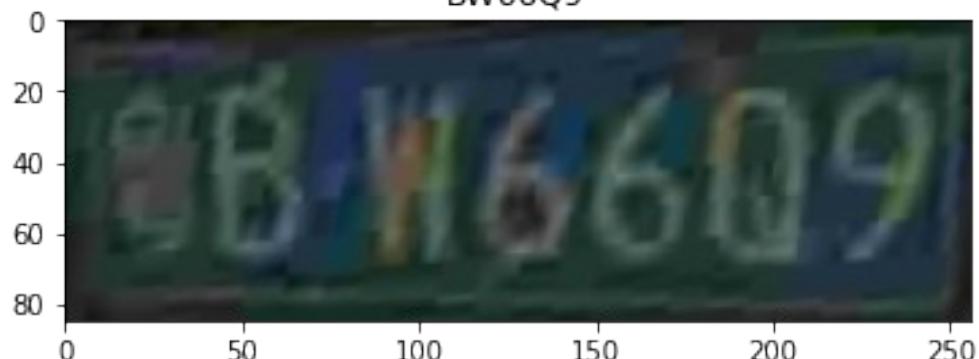
L57733



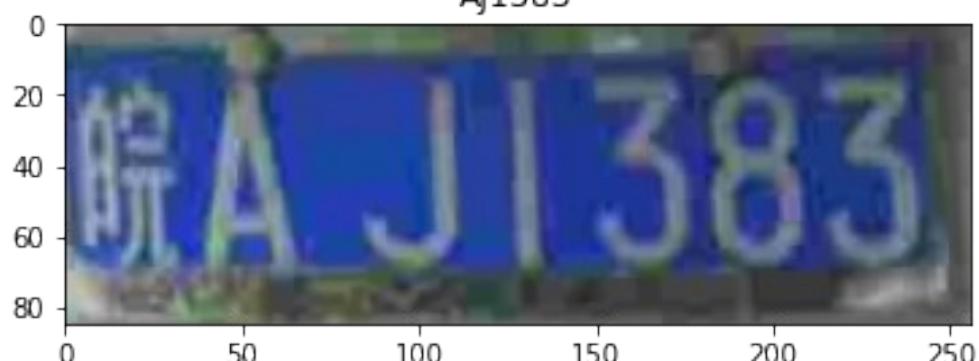
DLN338



BW66Q9



AJ1383



A379Y0



AY9Y29



A474B7



AQ1373



ALT223



AZ7Z57



AYR537



AL8Y46



AT278K



AJ300F



AWN293



AEW537



AUA010



AHS512



ACQ091



AY5T93



A69H67



HD798Z



A10R21



AJ6202



AR067S



A58G48



A4N328



AAC819



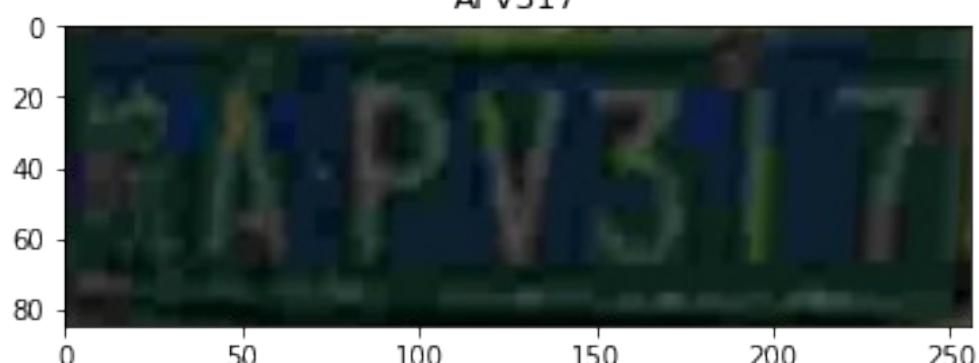
ABB303



AS8825



APV317



AN7K97



A7109F



A668L2



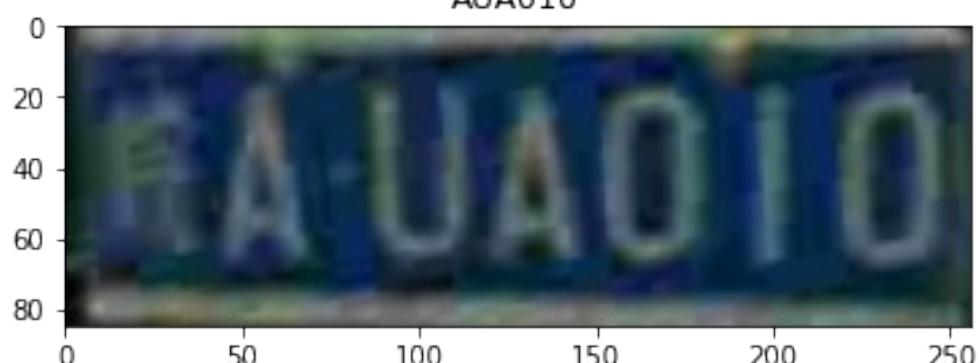
A61876



ACE819



AUA010



J92E33



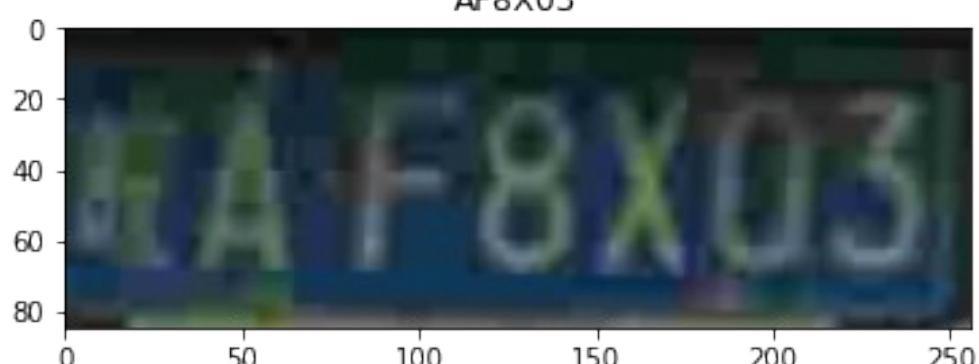
A5V645



AK0K58



AF8X03



ANA567



AEN444



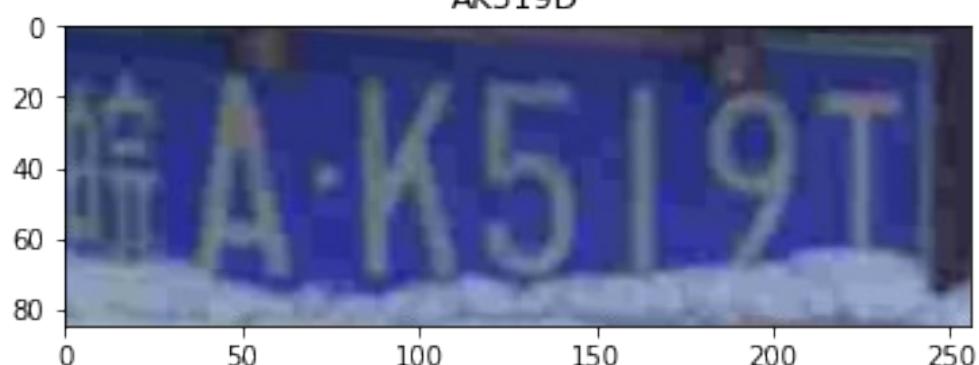
A1B521



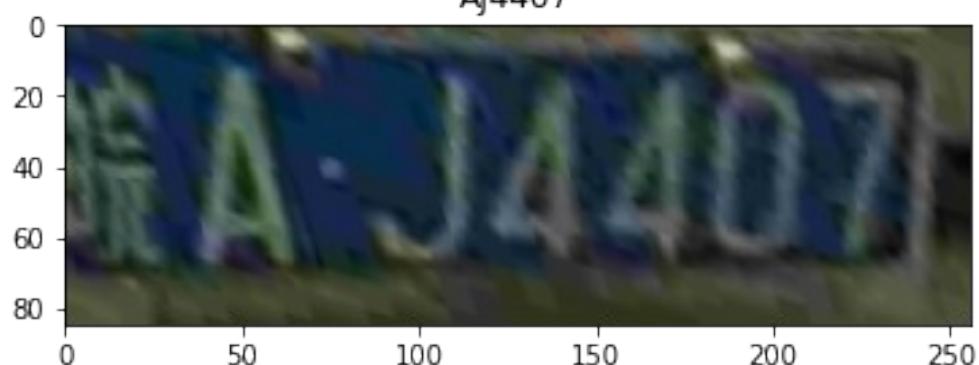
AZF762



AK519D



Aj4407



A781F7



A871M5



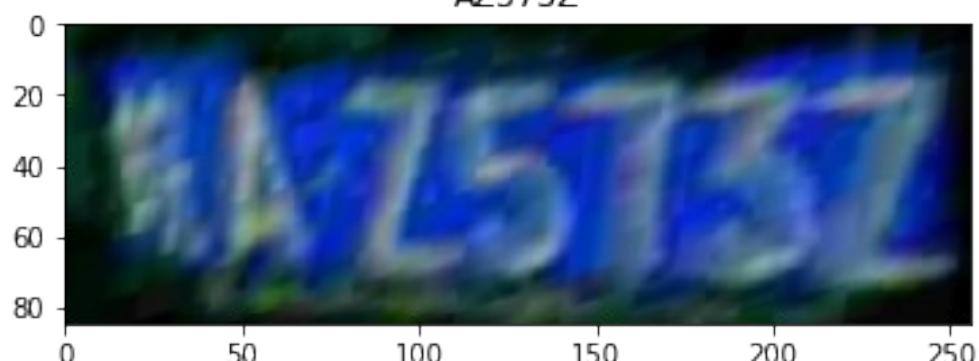
A685U0



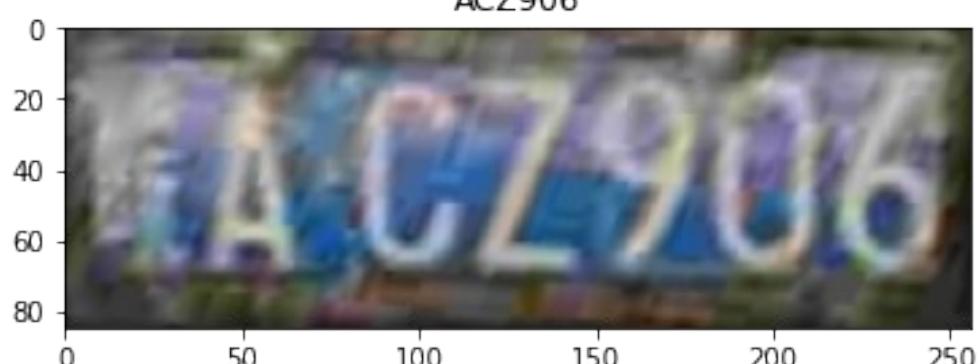
AV2939



AZ573Z



ACZ906



A947Q8



AD019T



E5692K



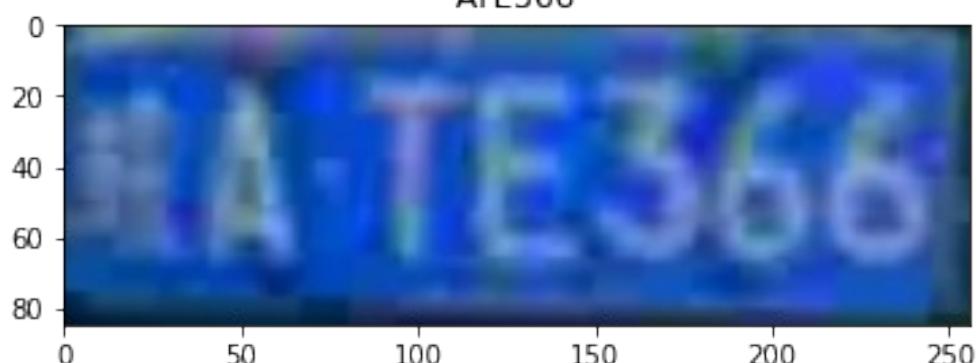
AUW299



A3Y735



ATE366



AK003H



HS1H77



E51QG7



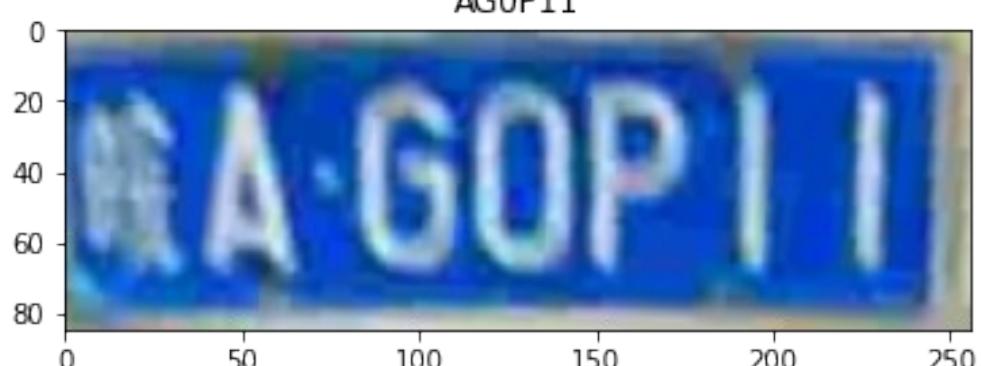
AC1V23



AR1732



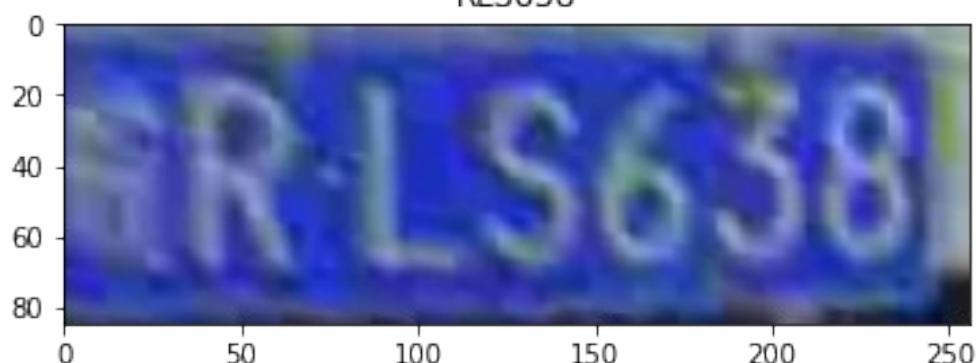
AG0P11



AJH019



RLS638



AOK195



AYB400



A551B1



A156L6



A81358



NQ0428



A48126



B54700



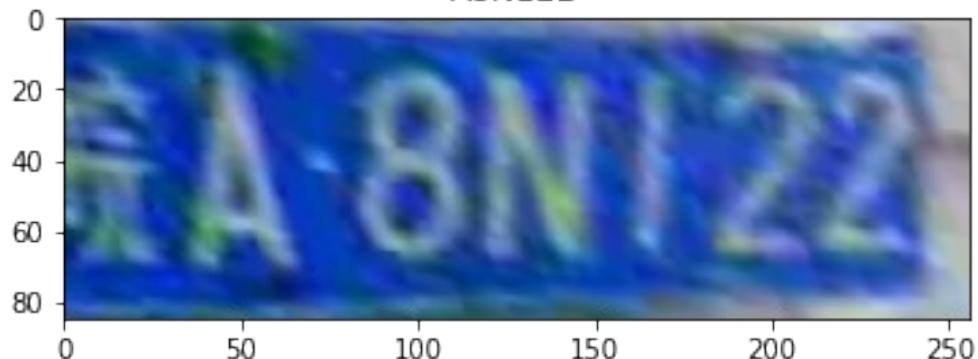
NPA667



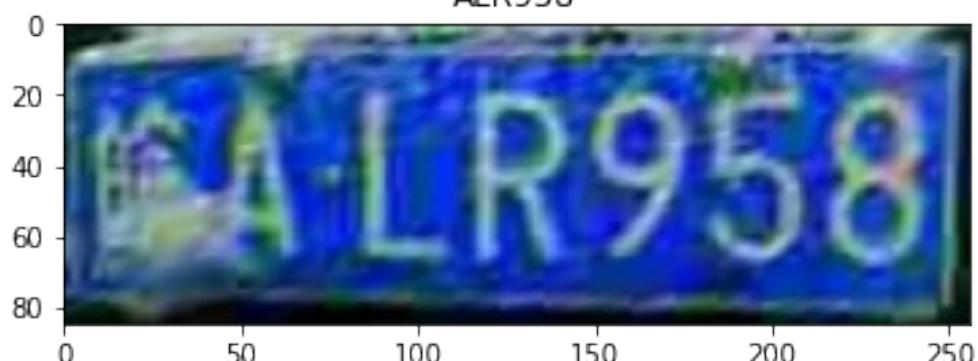
AC8638



A8N122



ALR958



AUE118



AUZ088

